

# The Traveling Salesman Problem (TSP)

## Overview

The Traveling Salesman Problem (TSP) is possibly *the* classic discrete optimization problem.

A preview :

- How is the TSP problem defined?
- What we know about the problem: NP-Completeness.
- The construction heuristics: Nearest-Neighbor, MST, Clarke-Wright, Christofides.
- K-OPT.
- Simulated annealing and Tabu search.
- The Held-Karp lower bound.
- Lin-Kernighan.
- Lin-Kernighan-Helsgaun.
- Exact methods using integer programming.



Our presentation will pull together material from various sources - see the references below. But most of it will come from [\[App12006\]](#), [\[John1997\]](#), [\[CS153\]](#).

## Defining the TSP


The TSP is fairly easy to describe:

- Input: a collection of points (representing cities).
- Goal: find a tour of minimal length.  
Length of tour = sum of inter-point distances along tour

**Input:**




**A non-optimal tour:**  
A B E D C



**The optimal tour:**  
A B C D E

- Details:
  - Input will be a list of  $n$  points, e.g.,  $(x_0, y_0), (x_1, y_1), \dots, (x_{n-1}, y_{n-1})$ .
  - Solution space: all possible tours.

- "Cost" of a tour: total length of tour.  
→ sum of distances between points along tour
- Goal: find the tour with minimal cost (length).
- Strictly speaking, we have defined the *Euclidean TSP*.
- There are really three kinds:
  - The Euclidean (points on the plane).
  - The *metric* TSP: triangle inequality is satisfied.
  - The *graph* TSP:



**Goal: find a minimal-length tour among tours that only use edges in the graph**

### Exercise:

- For an  $n$ -point problem, what is the size of the solution space (i.e., how many possible tours are there)?
- What's an example of an instance that's metric but not Euclidean?

Some assumptions and notation for the remainder:

- Let  $n = |V|$  = number of vertices.
- Euclidean version, unless otherwise stated.  
→ Complete graph.

## Some history

Early history:

- 1832: informal description of problem in German handbook for traveling salesmen.
- 1883 U.S. estimate: 200,000 traveling salesmen on the road
- 1850's onwards: circuit judges

**Exercise:** Find the following 14 cities in Illinois/Indiana on a map and identify the best tour you can:

Bloomington, Clinton, Danville, Decatur, Metamora, Monticello, Mt.Pulaski, Paris, Pekin, Shelbyville, Springfield, Sullivan, Taylorville, Urbana

- 1960's: Proctor and Gamble \$10K competition: a 33-city TSP.  
→ Won by a CMU mathematician (and others).
- A related problem: the *Knight's tour*.  
→ Start at bottom-left corner, and visit all squares exactly once and return to the start.

**Exercise:** Show how the Knight's tour can be converted into a TSP instance.

- The statisticians take an interest  
→ What is the expected length of an optimal tour for uniformly-generated points in 2D?
  - Several early analytic estimates in the 1940's.

- Famous Beardwood-Halton-Hammersley result [Bear1959]:  
If  $L^*$  = optimal tour's length then  $L^* / \sqrt{n} \rightarrow$  a constant  $\beta$
- $\beta$  estimated to be 0.72 for unit-square.
- Human solutions:
  - To assess problem-solving skill.
  - Part of some neurological tests.

TSP's importance in computer science:

- TSP has played a starring role in the development of algorithms.
- Used as a test case for almost every new (discrete) optimization algorithm:
  - Branch-and-bound.
  - Integer and mixed-integer algorithms.
  - Local search algorithms.
  - Simulated annealing, Tabu, genetic algorithms.
  - DNA computing.

Some milestones:

- Best known optimal algorithm: Held-Karp algorithm in 1962,  $O(n^2 2^n)$ .
- Proof of NP-completeness: Richard Karp in 1972 [Karp1972].  
→ Reduction from Vertex-Cover (which itself reduces from 3-SAT).
- Two directions for algorithm development:
  - Faster exact solution approaches (using linear programming).  
→ Largest problem solved optimally: 85,900-city problem (in 2006).
  - Effective heuristics.  
→ 1,904,711-city problem solved within 0.056% of optimal (in 2009)
- Optimal solutions take a long time  
→ A 7397-city problem took three years of CPU time.
- Theoretical development: (let  $L_H$  = tour-length produced by heuristic, and let  $L^*$  be the optimal tour-length)
  - 1976: Sahni-Gonzalez result [Sahn1976]. Unless  $P=NP$  no polynomial-time TSP heuristic can guarantee  $L_H/L^* \leq 2^{p(n)}$  for any fixed polynomial  $p(n)$ .
  - Various bounds on particular heuristics (see below).
  - 1992: Arora et al result [Aror1992]. Unless  $P=NP$ , there exists  $\epsilon > 0$  such that no polynomial-time TSP heuristic can guarantee  $L_H/L^* \leq 1 + \epsilon$  for all instances satisfying the triangle inequality.
  - 1998: Arora result [Aror1998]. For Euclidean TSP, there is an algorithm that is polynomial for fixed  $\epsilon > 0$  such that  $L_H/L^* \leq 1 + \epsilon$

## Approximate solutions: nearest neighbor algorithm

Nearest-neighbor heuristic:


- Possibly the simplest to implement.
- Sometimes called Greedy in the literature.
- Algorithm:

```

1.  V = {1, ..., n-1}           // Vertices except for 0.
2.  U = {0}                     // Vertex 0.
3.  while V not empty
4.      u = most recently added vertex to U
5.      Find vertex v in V closest to u
6.      Add v to U and remove v from V.
7.  endwhile
8.  Output vertices in the order they were added to U

```

**Exercise:** What is the solution produced by Nearest-Neighbor for the following 4-point Euclidean TSP. Is it optimal?




- What we know about Nearest-Neighbor:
  - $L_H/L^* \leq O(\log n)$
  - There are instances for which  $L_H/L^* = O(\log n)$
  - There are sub-classes of instances for which Nearest-Neighbor consistently produces the worst tour [Guti2007].


## Approximate solutions: the Clarke-Wright heuristic

The Clarke-Wright algorithm: [Clar1964].

- The idea:
  - First identify a "hub" vertex:




- Compute starting cost as cost of going through hub:



Starting cost = cost of visiting all points by walking  
along hub edges

- Identify "savings" for each pair of vertices:



shortcut from i to j

$$\text{savings}(i,j) = a + b - c$$

- Take shortcuts and add them to final tour, as long as no cycles are created.

- Algorithm:


```

1. Identify a hub vertex h
2.  $V_H = V - \{h\}$ 
3. for each  $i, j \neq h$ 
4.   compute  $\text{savings}(i, j)$ 
5. endfor
6. sortlist = Sort vertex pairs in decreasing order of savings
7. while  $|V_H| > 2$ 
8.   try vertex pair  $(i, j)$  in sortlist order
9.   if  $(i, j)$  shortcut does not create a cycle
10.    and  $\text{degree}(v) \leq 2$  for all  $v$ 
11.    add  $(i, j)$  segment to partial tour
12.    if  $\text{degree}(i) = 2$ 
13.       $V_H = V_H - \{i\}$ 
14.    endif
15.    if  $\text{degree}(j) = 2$ 
16.       $V_H = V_H - \{j\}$ 
17.    endif
18.  endif
19. endwhile
20. Stitch together remaining two vertices and hub into final tour


```

- Example (from above):

- Suppose vertex 4 is the hub vertex:




- Suppose (2,3) provides the most savings:




(2,3) gets added but no hub edge is removed (yet)

- Next, (1,2) gets added  
 →  $\text{degree}(2) = 2$   
 → must remove hub edge (2,4)




(1,2) is added and hub edge (2,4) is removed

- Continuing ... let's say we obtain:



Only two vertices, 5 and 7, remain connected to the hub

- Finally, add last two vertices and hub into final tour:



Final tour


- What's known about the CW heuristic:
  - Bound is logarithmic:  $L_H/L^* \leq O(\log n)$
  - Worst examples known:  $L_H/L^* \geq O(\log(n) / \log\log(n))$

## Approximate solutions: the MST heuristic


An approximation algorithm for (Euclidean) TSP that uses the MST: [\[Rose1977\]](#).

- The algorithm:
  1. First find the minimum spanning tree (using any MST algorithm).
  2. Pick any vertex to be the root of the tree.
  3. Traverse the tree in *pre-order*.
  4. Return the order of vertices visited in pre-order.


**Exercise:** What is the pre-order for this tree (starting at A)?



- Example:
  - Consider these 7 points:




- A minimum-spanning tree:




**Minimum spanning tree**

- Pick vertex A as the root:




- Traverse in pre-order:




**Visit order: A E B D C F G**

- Tour:




- Claim: the tour's length is no worse than twice the optimal tour's length.
  - Let  $L$  = the length of the tour produced by the algorithm.
  - Let  $L^*$  = the length of an optimal tour.
  - Let  $M$  = weight of the MST (total length).
  - Observe: if we remove any one edge from a tour, we will get a spanning tree.
    - $L^* > M$ .
  - Now consider a pre-order *tree walk* from the root, back to the root:





- Let  $W$  = length of this walk.
- Then,  $W = 2M$  (each edge is traversed twice).
- Thus,  $W < 2L^*$ .
- Finally, we will show that  $L \leq W$  and therefore,  $L < 2L^*$ .
- To see why, consider the tree walk from B to D:



$$\text{length BE} + \text{length ED} < \text{length BD}$$

→  $L$  takes a shorter route than  $W$  (triangle inequality).

- Thus,  $L \leq W$ .


What we know about this algorithm:

- The first heuristic to produce solutions within a constant of optimal.
- Easy to implement (since MST can be found efficiently).


## Approximate solutions: the Christofides heuristic

The Christofides algorithm: [Chri1976].

- First, as background, we need to understand two things:
  - What is an Euler tour (for general graphs)?
    - A tour that traverses all edges exactly once (but may repeat vertices)




*Euler tour exists*




*No Euler tour possible*

- Famous result: a graph has an Euler tour if and only if all its vertices have even degree.
- What is a minimal matching for a given subset of vertices  $V'$ ?
  - A "best" (minimal weight) subset of edges with the property that no edges have a common vertex




$V'$  = given set of vertices  
for which matching is desired




a matching

- Important result: min-matching can be found in poly-time.
- The key ideas in the algorithm:
  - First find the MST




- Then identify the odd-degree vertices



- There are an *even* number of such odd-degree vertices.

**Exercise:** Why?


- Find a minimal matching of these odd-degree vertices and add those edges



Adding "match" edges to original graph may result in multiple edges between a vertex pair


Drawn as curved to distinguish from regular graph edges

- Now all vertices have even degree.
- Next, find an Euler tour.



An Euler tour that may revisit vertices


- Now, walk along in Euler tour, but skip visited nodes



- This produces a TSP tour.

An improved bound:

- We will show that  $L_H/L^* \leq 1.5$
- Let  $M = \text{cost of MST}$ .  
→  $L^* \geq M$  (as argued before).
- Note: we are performing a matching on an even number of vertices.
- Now consider the original odd-degree vertices



- Consider the optimal tour on just these (even # of) vertices.
- Let  $L_O = \text{cost of this tour}$ .
- Let  $e^1, e^2, \dots, e^{2k}$  be the edges.
- Note:  $E_1 = \{e^1, e^3, \dots, e^{2k-1}\}$  is a matching.
- So is  $E_2 = \{e^2, e^4, \dots, e^{2k}\}$
- Now at least one set has weight at most  $L_O/2$ .  
→ Because both must add up to  $L_O$ .
- Also the optimal matching found earlier has less weight than either of these edge sets.  
→  $\text{min-match-cost} \leq L_O/2 \leq L^*/2$ .
- Thus  $\text{min-match-cost} + M \leq L^* + L^*/2$
- But  $L_H$  uses edges (or shortcuts) from min-match and MST  
→  $L_H \leq L^* + L^*/2$


Running time:

- Dominated by  $O(n^3)$  time for matching.
- Best known matching algorithm:  $O(n^{2.376})$

## K-OPT

Constructive vs. local-search heuristics:

- All four heuristics above were *constructive*  
→ A tour was built up step by step.
- In contrast, a local-search heuristic works as follows:




```

1.  T = some starting tour           // Perhaps by using Christofides.
2.  noChange = true
3.  repeat
4.    T' = makeChangeTo (T)
5.    if T' < T
6.      T = T'
7.      noChange = false
8.    endif
9.  until noChange
10. return T

```

## 2-OPT:

- Idea: Replace 2 edges and see if the cost improves.



- Find two edges and their endpoints.
- Swap endpoints.

- 2-OPT heuristic

```

1.  T = some starting tour
2.  noChange = true
3.  repeat
4.    for all possible edge-pairs in T
5.      T' = tour by swapping end points in edge-pair
6.      if T' < T
7.        T = T'
8.      noChange = false

```

```

9.         break           // Quit loop as soon as an improvement is found
10.      endif
11.  endfor
12. until noChange
13. return T

```

- An alternative: find best tour with all possible swaps:

```

1. T = some starting tour
2. noChange = true
3. repeat
4.   Tbest = T
5.   for all possible edge-pairs in T
6.     T' = tour by swapping end points in edge-pair
7.     if T' < Tbest
8.       Tbest = T'
9.     noChange = false
10.  endif
11. endfor
12. T = Tbest
13. until noChange
14. return T

```

K-OPT:

- 3-OPT is what you can get by considering replacing 3 edges.
- K-OPT considers K edges.
- Each K-OPT can be time-consuming for  $K > 3$ .

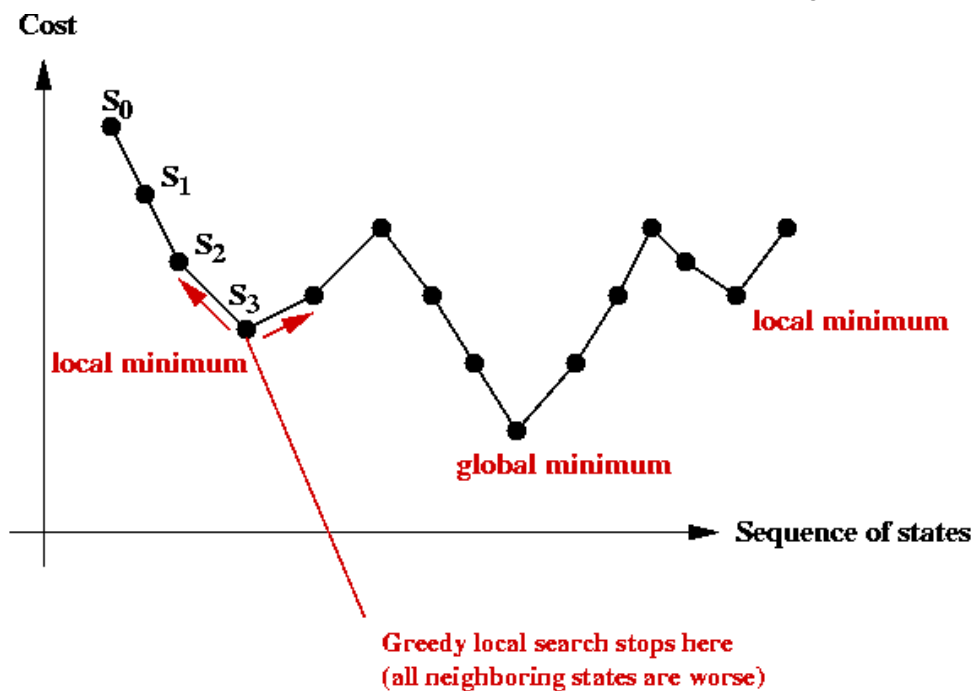
What we know about K-OPT:

- For general graphs:  $L_H/L^* \leq 0.25 n^{1/2k}$ .
- For Euclidean case,  $L_H/L^* \leq O(\log n)$ .
- In practice: 2-OPT and 3-OPT are much better than the construction heuristics.
- Note: Any K-OPT move can be reduced to a sequence of 2-OPT moves.  
→ But might it might require a long such sequence.


## Local Optima and Problem Landscape

Local optima:

- Recall: greedy-local-search generates one state (tour) after another until no better neighbor can be found  
→ does this mean the last one is optimal?
- Observe the trajectory of states:




- There is no guarantee that a greedy local search can find the (global) minimum.
- The last state found by greedy-local-search is a *local minimum*.  
→ it is the "best" in its neighborhood.
- The *global minimum* is what we seek: the least-cost solution overall.
- The particular local minimum found by greedy-local-search depends on the start state:



Problem landscape:


- Consider TSP using a particular local-search algorithm:
  - Suppose we use a graph where the vertices represent states.
  - An edge is placed between two "neighbors"  
e.g., for a 5-point TSP the neighbors of [0 1 2 3 4] are:




Neighbors of [0 1 2 3 4] using a 2-point swap

- The cost of each tour is represented as the "weight" of each vertex.
- Thus, a local-search algorithm "wanders" around this graph.
- Picture a 3D surface representing the cost *above* the graph.  
→ this is the problem landscape for a particular problem and local-search algorithm.

### STATES and COSTS



### LANDSCAPE



- A large part of the difficulty in solving combinatorial optimization problems is the "weirdness" in landscapes  
→ landscapes often have very little structure to exploit.
- Unlike continuous optimization problems, local shape in the landscape does NOT help point towards the global minimum.


Climbing out of local minima:

- A local-search algorithm gets "stuck" in a local minimum.
- One approach: re-run local-search many times with different starting points.
- Another approach (next): help a local-search algorithm "climb" out of local minima.


## Tabu search

Key ideas: [Glov1990].

- Suppose we decide to climb out of local minima.
- Danger: could immediately return to same local minima.



- In tabu-search, you maintain a list of "tabu tours".  
→ The algorithm avoids these.
- Each time you pick a minimum in a neighborhood, add that to the tabu list.




- Various alternatives to tabu-lists
  - Always add all neighborhood minimums.
  - Only add local minima.
- This way, Tabu forces more searching.
- A problem: a tabu-list can grow very long.
  - Need a *policy* for removing items, e.g.,
  - Least-recently used.
  - Throw out high-cost tours.

## Simulated annealing


Background:

- What is *annealing*?
  - *Annealing* is a metallurgic process for improving the strength of metals.
  - Key idea: cool metal slowly during the forging process.
- Example: making bar magnets:
  - Wrong way to make a magnet:
    1. Heat metal bar to high temperature in magnetic field.






2. Cool rapidly (quench):



◦ Right way: cool slowly (anneal)

• Why slow-cooling works:

◦ At high heat, magnetic dipoles are agitated and move around:



◦ The magnetic field tries to force alignment:



◦ If cooled rapidly, alignments tend to be less than optimal (local alignments):



◦ With slow-cooling, alignments are closer to optimal (global alignment):




• Summary: slow-cooling helps because it gives molecules more time to "settle" into a globally optimal configuration.

• Relation between "energy" and "optimality"

- The more aligned, the lower the system "energy".
- If the dipoles are not aligned, some dipoles' fields will conflict with others.
- If we (loosely) associate this "wasted" conflicting-fields with energy  
→ better alignment is equivalent to lower energy.
- Global minimum = lowest-energy state.

- The Boltzmann Distribution:
  - Consider a gas-molecule system (chamber with gas molecules):




- The state of the system is the particular snapshot (positions of molecules) at any time.
- There are high-energy states:



**High-energy: molecules bunched up**

and low-energy states:



**Low energy state: molecules spread apart**

- Suppose the states  $s_1, s_2, \dots$  have energies  $E(s_1), E(s_2), \dots$
- A particular energy value  $E$  occurs with probability

$$P[E] = Z e^{-E/kT}$$

where  $Z$  and  $k$  are constants.

- Low-energy states are more probable at low temperatures:
  - Consider states  $s_1$  and  $s_2$  with energies  $E(s_2) > E(s_1)$
  - The ratio of probabilities for these two states is:

$$r = P[E(s_1)] / P[E(s_2)] = e^{[E(s_2) - E(s_1)] / kT} = \exp ([E(s_2) - E(s_1)] / kT)$$

**Exercise :** Consider the ratio of probabilities above:

- Question: what happens to  $r$  as  $T$  increases to infinity?
- Question: what happens to  $r$  as  $T$  decreases to zero?

What are the implications?


Key ideas in simulated annealing: [\[Kirk1983\]](#).

- Simulated annealing = a modified local-search.
- Use it to solve a combinatorial optimization problem.
- Associate "energy" with "cost".
  - Goal: find lowest-energy state.
- Recall problem with local-search: gets stuck at local minimum.
- Simulated annealing will allow jumps to higher-cost states.
- If randomly-selected neighbor has lower-cost, jump to it (like local-search does).
- If randomly-selected neighbor is of higher-cost
  - flip a coin to decide whether to jump to higher-cost state

- Suppose current state is  $s$  with cost  $C(s)$ .
- Suppose randomly-selected neighbor is  $s'$  with cost  $C(s') > C(s)$ .
- Then, jump to it with probability

$$e^{-[C(s') - C(s)] / kT}$$

- Decrease coin-flip probability as time goes on:  
→ by decreasing temperature  $T$ .
- Probability of jumping to higher-cost state depends on cost-difference:



#### Implementation:

- Pseudocode: (for TSP)

**Algorithm:** TSPSimulatedAnnealing (points)

**Input:** array of points

```

1. // Start with any tour, e.g., in input order
   s = initial tour 0,1,...,n-1

   // Record initial tour as best so far.
2. min = cost (s)
3. minTour = s

   // Pick an initial temperature to allow "mobility"
4. T = selectInitialTemperature()

   // Iterate "long enough"
5. for i=1 to large-enough-number
   // Randomly select a neighboring state.
6.   s' = randomNextState (s)
   // If it's better, then jump to it.
7.   if cost(s') < cost(s)
8.     s = s'
   // Record best so far:
9.   if cost(s') < min
10.    min = cost(s')
11.    minTour = s'
12.   endif
13.  else if expCoinFlip (s, s')
   // Jump to s' even if it's worse.
14.    s = s'
15.  endif // Else stay in current state.
   // Decrease temperature.

```

```

16.         T = newTemperature (T)
17.     endfor

18.     return minTour

```

**Output:** best tour found by algorithm

**Algorithm:** randomNextState (s)

**Input:** a tour s, an array of integers

```
// ... Swap a random pair of points ...
```

**Output:** a tour

**Algorithm:** expCoinFlip (s, s')

**Input:** two states s and s'

```

1.  p = exp ( -(cost(s') - cost(s)) / T)
2.  u = uniformRandom (0, 1)
3.  if u < p
4.      return true
5.  else
6.      return false

```

**Output:** true (if coinFlip resulted in heads) or false

- Implementation for other problems, e.g., BPP
  - The only thing that needs to change: define a `nextState` method for each new problem.
  - Also, some experimentation will be need for the temperature schedule.

Temperature issues:

- Initial temperature:
  - Need to pick an initial temperature that will accept large cost increases (initially).
  - One way:
    - Guess what the large cost increase might be.
    - Pick initial  $T$  to make the probability 0.95 (close to 1).
- Decreasing the temperature:
  - We need a *temperature schedule*.
  - Several standard approaches:
    - Multiplicative decrease: Use  $T = a * T$ , where  $a$  is a constant like 0.99.  
→  $T_n = a^n$ .
    - Additive decrease: Use  $T = T - a$ , where  $a$  is a constant like 0.0001.
    - Inverse-log decrease: Use  $T = a / \log(n)$ .
  - In practice: need to experiment with different temperature schedules for a particular problem.

Analysis:

- How long do we run simulated annealing?
  - Typically, if the temperature is becomes very, very small there's no point in further execution  
→ because probability of escaping a local minimum is miniscule.
- Unlike previous algorithms, there is no fixed running time.
- What can we say theoretically?
  - If the inverse-log schedule is used  
→ Can prove "probabilistic convergence to global minimum"  
→ Loosely, as the number of iterations increase, the probability of finding the global minimum tends to 1.

In practice:

- Advantages of simulated annealing:
  - Simple to implement.
  - Does not need much insight into problem structure.
  - Can produce reasonable solutions.
  - If greedy does well, so will annealing.
- Disadvantages:
  - Poor temperature schedule can prevent sufficient exploration of state space.
  - Can require some experimentation before getting it to work well.
- Precautions:
  - Always re-run with several (wildly) different starting solutions.
  - Always experiment with different temperature schedules.
  - Always pick an initial temperature to ensure high probability of accepting a high-cost jump.
  - If possible, try different neighborhood functions.
- *Warning:*
  - Just because it has an appealing origin, simulated annealing is not guaranteed to work  
→ when it works, it's because it explores more of the state space than a greedy-local-search.
  - Simply running greedy-local-search on multiple starting points may be just as effective, and should be experimented with.

Variations:


- Use greedyNextState instead of the nextState function above.
  - Advantage: guaranteed to find local minima.
  - Disadvantage: may be difficult or impossible to climb out of a particular local minimum:
    - Suppose we are stuck at state  $s$ , a local minimum.
    - We probabilistically jump to  $s'$ , a higher-cost state.
    - When in  $s'$ , we will very likely jump back to  $s$  (unless a better state lies on the "other side").
  - Selecting a random next-state is more amenable to exploration.  
→ but it may not find local minima easily.
- Hybrid nextState functions:
  - Instead of considering the entire neighborhood of 2-swaps, examine some fraction of the neighborhood.
  - Switch between different neighborhood functions during iteration.
- Maintain "tabu" lists:
  - To avoid jumping to states already seen before, maintain a list of "already-visited" states and exclude these from each neighborhood.
- Thermal cycling:
  - Periodically raise temperature and perform "re-starts".
  - The idea is to force more exploration of the state space.

## The Held-Karp lower bound


Our presentation will follow the one in [\[Vale1997\]](#).

First, a definition:


- Consider a graph with vertices  $\{1, \dots, n\}$ :



- A 1-tree is a subgraph constructed as follows:




- Temporarily remove vertex 1 (and its edges) and find a spanning tree for vertices  $\{2, \dots, n\}$ .
- Then pick add two cheapest edges from vertex 1.
- Note: every tour (including the optimal one) is a 1-tree.



- The *min-1-tree* is the lowest weighted 1-tree among all 1-trees.  
→ This will be a lower bound for the optimal tour.
- A simple algorithm for the min-1-tree:
  - Find the MST for the graph without vertex 1.
  - Add the two cheapest edges from vertex 1.
- Is the min-1-tree a good bound?

**Exercise:** What is the difference between the optimal tour and the min-1-tree for this graph?




- The problem is: the MST can avoid using edges that the tour must take.

Held-Karp's idea:

- We will associate a  $\pi_i$ , a *vertex weight* with every vertex  $i$ .
- Define a modified graph  $G'$  as follows:
  - $G'$  has the same vertices and edges as  $G$ .
  - Let  $e_{ij}$  = weight of edge  $(i,j)$  in  $G$ .
  - Let  $c_{ij}$  = weight of edge  $(i,j)$  in  $G'$ .
  - Then define  $c_{ij} = e_{ij} + \pi_i + \pi_j$ .

- For example:



**Exercise:** What is the difference between the min-1-tree and the optimal tour for the above modified graph  $G'$ ? What vertex weight for the top-right vertex best closes the gap between the min-1-tree and the optimal tour?

- Thus, one can *choose* weights so that the min-1-tree is as high as possible in  $G'$ .

In more detail:

- Let  $T$  be a 1-tree and  $T'$  be a tour.
- Let  $d_i^T$  = the degree of node  $i$  in  $T$ .
- Let  $L(T, G) = \text{cost of 1-tree } T \text{ using graph } G$ .
- Let  $L(T', G) = \text{cost of tour } T' \text{ using graph } G$ .
- Since every tour is a 1-tree,  $\min_T L(T, G) \leq \min_{T'} L(T', G)$ .
- Now, for a 1-tree  $T$ ,  

$$L(T, G') = L(T, G) + \sum_{i \in T} (d_i^T - 2)\pi_i.$$
- Similarly, for a tour  $T'$ ,  


$$L(T', G') = L(T', G) + \sum_{i \in T'} 2\pi_i.$$
- Thus, subtracting and taking minimum,  $\min_T L(T, G) + \sum_{i \in T} (d_i^T - 2)\pi_i \leq \min_{T'} L(T', G) = L^*$  (the optimal tour).
- To summarize, we want to find the min-1-tree with weights  $\pi$  and then correct for that by subtracting off the additional weights.
- Let  $W(\pi) = \min_T L(T, G) + \sum_{i \in T} (d_i^T - 2)\pi_i$ .
- Then, the desired "best" Held-Karp bound is:  $\max_{\pi} W(\pi)$ .

An optimization procedure:

- Let  $V_{T(\pi)}$  be the vector  $(d_1^T, \dots, d_n^T)$ .
- Let  $C_{T(\pi)}$  be the cost of min-1-tree using  $\pi$ .
- Then, write  $W(\pi) = C_{T(\pi)} + \pi \cdot V_{T(\pi)}$ .
- Next, suppose that  $\pi'$  is a vector in  $\pi$ -space such that  $W(\pi') \geq W(\pi)$ .
- Then, Held-Karp show that  $(\pi' - \pi) \cdot V_{T(\pi)} \geq 0$ .
- This means that larger values of  $W(\pi')$  are in the right half-space pointed to by the vector  $V_{T(\pi)}$ .
- Next step: an iterative optimization procedure.


First, a little background on gradient-based optimization:

- Consider a (single-dimensional) function  $f(x)$ :



- Let  $f'(x)$  denote the derivative of  $f(x)$ .
- The gradient at a point  $x$  is the value of  $f'(x)$ .
  - Graphically, the slope of the tangent to the curve at  $x$ .

- Observe the following:



- To the left of the optimal value  $x^*$ , the gradient is negative.
- To the right, it's positive.

- We seek an iterative algorithm of the form

```
while not over
  if gradient < 0
    move rightwards
  else if gradient > 0
    move leftwards
  else
    stop // gradient = 0 (unlikely in practice, of course)
  endif
endwhile
```

- The gradient descent algorithm is exactly this idea:

```
while not over
   $x = x - \alpha f'(x)$ 
endwhile
```



Here, we add a scaling factor  $\alpha$  in case  $f(x)$  values are of a different order-of-magnitude:



Back to vertex-weight optimization:

- Unfortunately, we don't have a differentiable function.
- For this case, the Russian mathematician Polyak devised what's called the *sub-gradient* algorithm:
  - For a differentiable function, the gradient "points" in the right direction.
  - For a non-differentiable function, it's still possible to use a gradient that points in the right direction.
- For the vertex-weights, the iteration turns out to be:  $\pi_i^{(m+1)} = \pi_i^{(m)} + \alpha^{(m)} (d_i - 2)$ .
- Intuitively, this means:
  - Increase the weights for vertices with 1-min-tree degree  $> 2$ .
  - Decrease the weights for vertices with 1-min-tree degree  $< 2$ .
  - Thus, the iteration tries to force the 1-min-tree to be "tour-like".
- Polyak showed that sub-gradient iteration works if the stepsizes  $\alpha^{(m)}$  are chosen properly:
  - $\alpha^{(m)} \rightarrow 0$
  - $\sum_m \alpha^{(m)} = \infty$
- To summarize:
  - Start with some vector of vertex-weights  $\pi$ .
  - Repeatedly apply the iteration  $\pi_i^{(m+1)} = \pi_i^{(m)} + \text{stepsize} * \text{sub-gradient } V_{T(\pi)}$ .
- Implementation issues:
  - Each iteration requires an MST computation.
    - Can be expensive for large  $n$ .
  - One approximation: reduce number of edges by considering only best  $k$  neighbors (e.g.,  $k=20$ ).

## The Lin-Kernighan algorithm


Key ideas:

- Devised in 1973 by Shen Lin (co-author on BB(N) numbers) and Brian Kernighan (the "K" of K&R fame).
- Champion TSP heuristic 1973-89.
- LK is iterative:
  - Starts with a tour and repeatedly improves, until no improvement can be found.
- Idea 1: Make the  $K$  edges in  $K\text{-OPT}$  *contiguous*




- This is just the high-level idea
  - The algorithm actually alternates between a "current-tour-edge" and a "new-putative-edge".
- Let the K in K-OPT vary at each iteration.
  - Try to increase K gradually at each iteration.
  - Pick the best K (the best tour) along the way.
- Allow some limited backtracking.
- Use a tabu-list to create freshness in exploration.

Note: we will use an artificial depiction of a tour as follows:



**Actual graph and tour**




**Circular visualization of tour**


This will be used to explain some ideas.

The LK algorithm in more detail:


- At each iteration, LK identifies a sequence of edges  $x_1, y_1, x_2, y_2, \dots, x_k, y_k$  such that:




- Each  $x_i$  is an edge in the current tour.
- Each  $y_i$  is NOT in the current tour.
- They are all unique (no repetitions).
- The last  $y_k$  returns to the starting point  $t_1$
- We'll call this an *LK-move*.
- For example:




- Notice that if we stop at any intermediate  $y_i$ , we get a 1-tree.



- Let  $G_1$  = gain after first x-y-pair:  
 $G_1 = c(x_1) - c(y_1)$
- Similarly,  
 $G_2 = c(x_1) - c(y_1) + c(x_2) - c(y_2)$ .
- Gain criterion used by algorithm:  
 Keep increasing  $k$  as long as  $G_k > 0$ .
- Note: this is a non-trivial addition because it allows for a temporary loss in gain:



- Neighbor limitation:



- LK limits the number of neighbors to the  $m$  nearest neighbors, where  $m$  is an algorithm parameter (e.g.,  $m=10$ ).
- Re-starts:
  - Recall: there are  $n$  choices for  $t_1$ , the very first node.
  - LK tries all  $n$  before giving up.
- Best-tour: at all times LK records the best tour found so far.
- Note: LK is actually a little more complicated than described above, but these are the key ideas.

#### Performance:

- The standard heuristics (construction, K-OPT) give tours with 2-5% above Held-Karp.
- LK is usually between 1-2% off.


## LKH-1: Lin-Kernighan-Helsgaun

From 1999-2009, Keld Helsgaun [Hels2009], added a number of sophisticated optimizations to the basic LK algorithm:

- The first set were added in 1999: [Hels1999].
  - We'll call this LKH-1.
- And the second set in 2009: [Hels2009].
  - We'll call this LKH-2.

#### Key ideas in LKH-1:

- Use  $K=5$  (prefer this value of  $K$  over smaller ones).
  - Experimental evidence showed that the improvement going from 4- to 5-OPT is much better than 3- to 4-OPT.
  - Tradeoff: if  $K$  is too high, it takes too long
    - Fewer iterations
    - Less exploration of search space (even if you search a particular neighborhood more thoroughly).
- Relax *sequentiality* allow some  $x_i$ 's and  $y_i$ 's to repeat.
- Replace closest  $m$  neighbors with a different set of  $M$  neighbors:
  - Problem with LK:




- Recall best 1-tree in Held-Karp bound?
  - Many of these edges are "good" edges for the tour.
  - Experimental evidence: 70-80% of these edges are in optimal tour.
- LKH-1 idea: prefer 1-tree edges that go to neighbors.
- Let  $L(T)$  = cost of best 1-tree
  - Can be computed fast (MST)
- For any edge  $e$ , let
  - $L(T,e)$  = cost of best 1-tree that *must* use  $e$ .
- How to force using an edge  $e$ ?
  - Find min-1-tree.
  - Add  $e$  to tree.
  - This causes a cycle.
  - Remove heaviest edge in cycle.
  - This leaves a min-1-tree that uses  $e$ .
- Define  $\alpha(e) = L(T,e) - L(T)$  = importance of  $e$  in "1-tree-ness"
- Note:  $\alpha(e)=0$  for any edge in min-1-tree.
- LKH-1 sorts neighbors by  $\alpha$  and uses best  $m$  of these.

## LKH-2: Lin-Kernighan-Helsgaun, Part 2

Key additions to LKH-1:

- Allow  $K$  to increase beyond 5.
- Problem-partitioning:




- Divide points into clusters.
- Find best tour for each cluster.
- Stitch together into final tour.
- Run algorithm many times and merge "best parts" from multiple tours.
  - Called *iterative partial transcription*.
- Use sophisticated tour data structures to speed up running time.
- Results: million city problem with 0.058% of Held-Karp.
  - Within 0.058% of optimal.

Let's examine the partitioning idea:

- LKH-2 tries a number of partitionings, using different clustering algorithms.
- K-means clustering:

1. **repeat**  
     // Note: this is a different K than in K-OPT.
2.     Pick k centroids.
3.     Assign each point to closest centroid.
4.     Re-compute the centroid based on assignments.
5. **until** no change

- Tour segmentation:
  - Run LKH-2 once to find a tour.
  - Segment the tour and re-solve the segments (partition).
- Geometric:









Recursive subdivision of space  
 (similar to k-d trees or quad-trees)

- Space-filling curve:

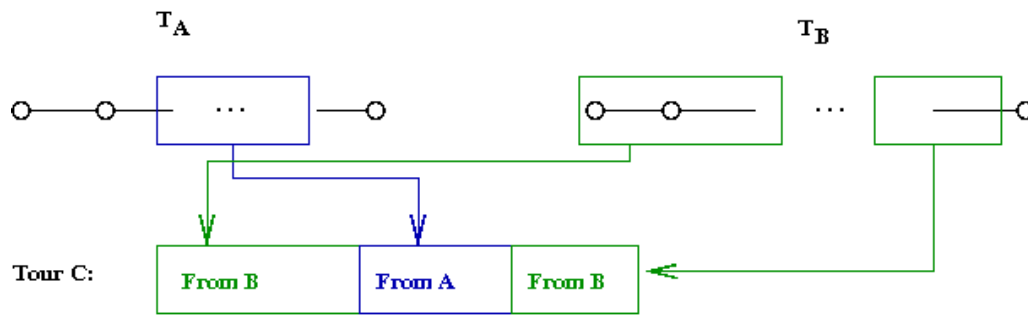
Hilbert space-filling curves (recursively defined)

Points close by along curve  
 are likely to be clustered




Iterative partial transcription (IPT):

- This is an idea from [\[Mobi1999\]](#).
- Goal: given two tours  $T_A$  and  $T_B$ , compute  $T_C$  that is better than both  $T_A$  and  $T_B$ .



- A single IPT *trial-swap* between tours  $T_A$  and  $T_B$  to creates tours  $T_{A'}$  and  $T_{B'}$




- An IPT-iteration:
  - Identifies all possible valid swap segments.
  - Tries the swaps and identifies the best possible tour that can be generated.
- How to use IPT:
  - Generate  $m$  tours  $T_1, \dots, T_m$ .
  - For each pair of tours  $i, j$ , perform an IPT-iteration.

## Data structures

Given a K-OPT move, is the resulting "tour" a valid tour?


- Example:



- Naive way: walk along new tour  $T'$  to see if all vertices are visited  
→  $O(n)$  per trial edge-swap
- Another problem: how to maintain tours?

Operations on tour data structures:

- First, note that any single swap can result in reversing the tour order for one of the segments affected:




- A single 2-OPT move will be called a *flip* operation.
- Also, any K-OPT move can be implemented by a sequence of 2-OPT moves.  
→ LK-MOVE can be written to use *flip* operations.
- Other operations that need to be supported:
  - $\text{next}(a)$ : the next node in tour order.
  - $\text{prev}(a)$ : the previous node in tour order.
  - $\text{between}(a,b,c)$ : determine whether  $b$  is between  $a$  and  $c$  in tour-order.
- Note: If a flip is performed correctly, it will result in a valid tour.
- Fredman et al. [Fred1995] show a lower bound of  $(\log n) / (\log \log n)$  for these operations.

Arrays:

- Simple to implement.




- But consider what needs to be done to reverse a segment:



→ Can take  $O(n)$ .

Doubly-linked lists:


- flip* takes  $O(1)$  pointer manipulations.
- Order reversal is also easy (comes for free):  $O(1)$ .




- But finding elements is hard:  $O(n)$ .

Modified splay trees:


- What is a splay tree?
  - Also called a *self-adjusting binary tree*.
  - [See lecture in algorithms course.](#)
  - Recall problem with binary trees: can go out of balance.
  - Problem with forced balance (e.g. AVL): too much overhead.
    - But use of *rotations* is useful.
  - Example of a splay-step: two mini-rotations:



- Another example:




- In a splay-tree: every accessed node is *splayed to the root*.  
→ Similar to Move-to-Front in linked lists.
- Using a splay-tree for a tour:
  - Each node represents a city.
  - Initially, for first tour: in-order traversal is the tour:




**Exercise:** What is the tour represented by the above tree?


- Reversals are noted by marking intermediate nodes, e.g.



- Each time a reversed-node is encountered, switch order (left swapped with right) in in-order traversal:




- Maintain an external array of pointers into tree, one per node.
- Implementing  $next(a)$ :
  - Recall  $next(a)$  in ordinary binary trees: leftmost node of the right subtree.



- Locate  $a$  using pointer-array:  $O(1)$ .
- Splay to root.
- Find successor using *tour-order* (instead of numeric order).
  - With no reversals, this is the leftmost node of the right subtree.
- With reversals, need to change direction for each flip (when recursing).
- The most complex operation is *flip()*:
  - Just like the splay-tree, there are several different cases.
  - Many involve some type of reversal.
  - The general idea (an example):

The flip(a,b,c,d) operation



The segment tree:

- Devised by Applegate and Cook.
- Based on key observation about LK:
  - You try a sequence of flips (the LK-move).
  - When it doesn't work, you discard the whole sequence.
- In the data structures so far:
  - Every flip changes the data structure.
  - To discard, we need to *undo* flips in reverse order.
- A segment-tree tries to avoid the *undo* part.
  - Array representation of tour.
  - An auxiliary segment-list:
    - To help with tentative flips.
  - An auxiliary segment tree:
    - To help with fast navigation.

Performance:

- Segment-tree is usually best.
- 2-level list is next.
- Splay tree next (with theoretically the best performance).

## Exact solution techniques: background

The general idea:

- Formulate TSP as a Integer Programming (IP) problem.
- Apply the *cutting-plane* approach.

- Judicious choice of cutting-plane heuristics.

But, first, what is Integer Programming? We'll need some background in *linear programming*.

Linear programming:

- The word *program* has different meaning than we are used to.  
→ More like a "programme" of events.
- An LP (Linear Programming) problem is (in standard form):

$$\begin{aligned}
 &\max \quad c_1x_1 + c_2x_2 + \dots + c_nx_n \\
 &\text{such that} \\
 &\quad a_{11}x_1 + \dots + a_{1n}x_n \leq b_1 \\
 &\quad a_{21}x_1 + \dots + a_{2n}x_n \leq b_2 \\
 &\quad \vdots \\
 &\quad a_{n1}x_1 + \dots + a_{nn}x_n \leq b_n \\
 &\text{and} \\
 &\quad x_i \geq 0, \quad i=1, \dots, n \\
 &\quad x_i \in \mathbb{R}
 \end{aligned}$$

- In vector/matrix notation:

$$\begin{aligned}
 &\max \quad c^T x \\
 &\text{s.t.} \quad Ax \leq b \\
 &\quad x \geq 0
 \end{aligned}$$

- Example:

3.6 = selling price per unit of product 1

Amount of product 2 =  $x_2$


$$\begin{aligned}
 &\max \quad 3.6x_1 + 14x_2 \\
 &\text{s.t.} \quad 12x_1 + 23.47x_2 \leq 400 \\
 &\quad 34x_1 + 16x_2 \leq 650 \\
 &\quad 55x_1 \leq 1200 \\
 &\text{and } x_i \geq 0
 \end{aligned}$$

only product 1 needs raw material #3


650 = total amount of raw material #2

16 = amount of raw material #2 used per unit of product 2


- Geometric intuition of inequality constraints ( $Ax \leq b$ ):



- Each inequality defines a half-plane (half-space).
  - The intersection is a polytope (polygon in 2D).
  - The feasible region is sometimes called the *simplex*.
- If we plot objective function "lines":



- If we make a line-equation out of the objective function, some lines will pass through the feasible region.
  - Clearly, we want the line with the highest "value" (for a max problem).
  - Sweeping the line upwards (higher value), we want the line that is the last line to intersect the feasible region.
  - This line always intersects the region at a *corner*.
- Three key algorithms, all major milestones in the development of LP:
    - George Dantzig's Simplex algorithm (1947).
    - Leonid Khachiyan's ellipsoid method (1979).
    - Narendra Karmarkar's interior-point method (1984).
  - The simplex method:



- Start at a corner in the feasible region.
  - A *simplex-move* is a move to a neighboring corner.
  - Pick a better neighbor to move to (or even best neighbor).
  - Repeat until you've reached optimal solution.
- What's known about the simplex method:
    - Guaranteed to find optimal solution.
    - Worst-case running time: exponential.
    - In practice, it's quite efficient, approximately  $O(n^3)$ .
    - Very efficient implementations available, both commercial and open-source.
    - Has been used to solve very large problems (thousands of variables).
  - What's known about the other algorithms:
    - Khachiyan's ellipsoid method: provably polynomial, but inefficient in practice.
    - Karmarkar's algorithm: provably polynomial and practically efficient for many types of LP problems.

- Note: an LP problem with equality constraints

$$\begin{array}{ll} \max & c^T x \\ \text{s.t.} & Ax = b \\ & x \geq 0 \end{array}$$

can be converted to an equivalent one in standard form (with inequality constraints).

- Similarly, a min-problem can be converted to a max-problem.

Integer programming:


- An integer program (IP) is an LP problem with one additional constraint: all  $x_i$ 's are required to be integer:

$$\begin{array}{ll} \max & c^T x \\ \text{s.t.} & Ax \leq b \\ & x \geq 0 \\ & x \in \mathbb{Z} \end{array}$$

## Exact solution techniques: TSP as an IP problem

First, let's express TSP as an IP problem:


- We'll assume the TSP is a Euclidean TSP (the formulation for a graph-TSP is similar).
- Let the variable  $x_{ij}$  represent the directed edge  $(i,j)$ .
- Let  $c_{ij} = c_{ji}$  = the cost of the undirected edge  $(i,j)$ .



- Consider the following IP problem:

$$\begin{array}{ll} \min & \sum_{i,j} c_{i,j} x_{i,j} \\ \text{s.t.} & \sum_j x_{i,j} = 1 \quad // \text{ Only one outgoing arc from } i \\ & \sum_i x_{i,j} = 1 \quad // \text{ Only one incoming arc at } j \end{array}$$


- Unfortunately, this is not sufficient:



You can get multiple cycles.

→ Called *sub-tours*

- What to do? Consider this idea:



- Consider a subset of vertices  $S$ .
- In a valid tour,
 
$$\sum_{i,j} x_{ij} \leq |S|-1 \text{ for all } i,j \in S.$$
- This is an inequality constraint that could be added to the IP problem.
 

→ Called a *sub-tour* constraint.
- How many such constraints need to be added to the IP problem?
 


→ One for each possible subset  $S$ .

→ Exponential number of constraints!
- Fortunately, one can add these constraints only as and when needed (see below).

Solving the IP problem:


- Naive approach:
    - Solve the *LP relaxation* problem first.
 

→ Remove integer constraints (temporarily) to get a regular LP, and solve it.
    - Round LP solution to nearest integers.
- Unfortunately, this may not yield a feasible solution:




- Branch-and-bound:
  - We'll explain this for 0-1-IP problems (variables are binary-valued).
  - First, consider a simple exhaustive search, organized as a tree-search (the "branch" part):






- The tree itself can be explored in a variety of ways:
  - Breadth-first (high memory)
    - High memory requirements.
  - Depth-first
    - Low memory requirements.
  - Cost-first
    - Expand the node that adds the least overall cost to the (partial) objective function.
- Note: if the cost to a node already exceeds the best tour so far, there's no need to explore further.
  - Parts of the tree can be *pruned*.



- Cutting planes:



- Add constraints to force the LP-solutions towards integers.
- With a sequence of such constraints, such a process can converge to an integer solution.
- However, it can take a long time.
- Gomory's algorithm:
  - A general cutting-plane algorithm for any IP.
  - The idea:
    - Solve LP.
    - Examine equations satisfied at corner point (of LP).
    - Round to integers in inequalities involving those variables.
    - Add these to constraints.
    - Repeat.

- Unfortunately, it is slow in practice.

### History of applying IP to TSP:


- Original cutting plane idea due to Dantzig, Fulkerson and Johnson in 1954.
  - Idea:

```

repeat
    solve LP
    identify sub-tours (cycles) and add corresponding " $|S|-1$ " constraints.
until full-tour found

```

- Dantzig et al added a few more "sub-tour" like constraints.
- Today, there are several families of cutting-plane constraints for the TSP.
- Branch-and-cut
  - Cutting planes "ruled" until 1972.
  - Saman Hong (JHU) in 1972 combined cutting-planes with branch-and-bound
    - Called branch-and-cut.
  - The idea: some variables might change too slowly with cutting planes
    - For these, try both 0 and 1 (branch-and-bound idea).
  - Alternate way of viewing this:



- More sophisticated "cut" families:
  - Grotschel & Padberg, 1970's.
  - Padberg and Hong, 1980: 318-city problem.
  - Grotschel and Holland, 1987: 666-city problem.
  - Padberg and Rinaldi, 1987-88: combined multiple types of cuts, branch-and-cut and various tricks to solve 2392-city problem.
- During this time, LP techniques improved greatly
  - Can cut down "active" variables in an LP problem.
- Applegate et al (2006)
  - Sophisticated LP techniques, new data structures.
  - 85,900 city problem.

---

## References and further reading

[WP-1] [Wikipedia entry on TSP.](#)

[WP-1] [Georgia Tech website on TSP.](#)

[Appl2006] D.L.Applegate, R.E.Bixby, V.Chvatal and W.J.Cook. *The Traveling Salesman Problem*, Princeton Univ. Press, 2006.

[Aror1992] S.Arora, C.Lund, R.Motwani, M.Sudan and M.Szegedy. Proof verification and hardness of approximation problems. *Proc. Symp. Foundations of Computer Science*, 1992, pp.14-23.

[Aror1998] S.Arora. [Polynomial Time Approximation Schemes for Euclidean Traveling Salesman and other Geometric Problems.](#) *JACM*, 45:5, 1998, pp. 753-782.

[Bear1959] J.Beardwood, J.H.Halton and J.M.Hammersley. The Shortest Path Through Many Points. *Proc. Cambridge Phil. Soc.*, 55, 1959, pp.299-327.

- [Chan1994] B.Chandra, H.Karloff and C.Tovey. [New results on the old k-opt algorithm for the TSP](#). *5th ACM-SIAM Symp. on Discrete Algorithms*, 1994, pp.150-159.
- [Clar1964] G.Clarke and J.W.Wright. Scheduling of vehicles from a central depot to a number of delivery points. *Op.Res.*, 12, 1964, pp.568-581.
- [Chri1976] N.Christofides. Worst-case analysis of a new heuristic for the travelling salesman problem. Report No. 388, GSIA, Carnegie-Mellon University, Pittsburgh, PA, 1976.
- [Croe1958] G.A.Croes. A method for solving traveling salesman problems. *Op.Res.*, 6, 1958, pp.791-812.
- [Fred1995] M.L.Fredman, D.S.Johnson, L.A.McGeogh and G.Ostheimer. Data structures for traveling salesmen. *J.Algorithms*, Vol.18, 1995, pp.432-479.
- [Glov1990] F.Glover. [Tabu Search: A Tutorial](#), *Interfaces*, 20:1, 1990, pp.74-94.
- [Guti2007] G.Gutin and A.Yeo. [The Greedy Algorithm for the Symmetric TSP](#). *Algorithmic Oper. Res.*, Vol.2, 2007, pp.33--36.
- [Held1970] M.Held and R.M.Karp. The traveling-salesman problem and minimum spanning trees. *Op.Res.*, 18, 1970, pp.1138-1162.
- [Hels1998] K. Helsgaun. [An Effective Implementation of the Lin-Kernighan Traveling Salesman Heuristic](#), DATALOGISKE SKRIFTER (Writings on Computer Science), No. 81, 1998, Roskilde University.
- [Hels2009] K. Helsgaun. [General k-opt submoves for the Lin-Kernighan TSP heuristic](#). *Mathematical Programming Computation*, 2009.
- [John1997] D.S.Johnson and L.A.McGeoch. [The Traveling Salesman Problem: A Case Study in Local Optimization](#). In Aarts, E. H. L.; Lenstra, J. K., *Local Search in Combinatorial Optimisation*, John Wiley and Sons Ltd, pp. 215-310, 1997.
- [Karp1972] R.Karp. [Reducibility among combinatorial problems](#), in R. E. Miller and J. W. Thatcher (editors). , New York: Plenum. pp. 85-103.
- [Kirk1983] S.Kirkpatrick, C.D.Gelatt, and M.P.Vecchi. Optimization by Simulated Annealing. *Science*, 220 1983, pp.671-680.
- [Lin1973] S.Lin and B.W.Kernighan. [An Effective Heuristic Algorithm for the Traveling- Salesman Problem](#). *Op.Res.*, 21, 1973, pp.498-516.
- [Mobi1999] A.Mobius, B.Freisleben, P.Merz and M.Schreiber. Combinatorial optimization by iterative partial transcription. *Phys.Rev. E*, 59:4, 1999, pp.4667-74.
- [Rein1991] G.Reinelt. [TSPLIB - A Traveling Salesman Problem Library](#). *ORSA J. Comp.*, 3:4, 1991, pp. 376-384.
- D.J.Rosenkrantz, R.E.Stearns and P.M.Lewis. An analysis of several heuristics for the traveling salesman problem. *SIAM J. Computing*, Vol.6, 1977, pp.563-581.
- [Sahn1976] S.Sahni and T.Gonzalez. P-complete approximation problems. *J.ACM*, Vol.23, 1976, pp.555-565.
- [CS153] R.Simha. [Course notes for CS-153](#) (Undergraduate algorithms course).
- [Vale1997] C.L.Valenzuela and A.J.Jones. [Estimating the Held-Karp lower bound for the geometric TSP](#). *European J. Op. Res.*, 102:1, 1997, pp.157-175.

Note: The Hilbert curve was an image found on Wiki-commons.

---